Epistemic Honesty in Predictive Systems: An Impossibility Result

Anonymous Author(s)

Abstract

Large language model "hallucinations" are commonly framed as engineering failures to be solved with better training, more data, or improved alignment techniques. We argue this framing is incorrect. We present a formal impossibility result demonstrating that *epistemic honesty*, which is the property that a system's outputs faithfully reflect its internal epistemic state, is unverifiable for systems operating under text-only observation models with bounded supervisors.

We decompose this impossibility into two theorems. **Theorem 1 (Representational Impossibility)** proves that policies conditioning solely on queries cannot satisfy epistemic requirements that depend on external world states. **Theorem 4.8 (Learnability Impossibility)** proves that even architectures with retrieval capabilities cannot learn honest policies via preference optimization when supervisors are resource-bounded, because plausible fabrications are observationally indistinguishable from truthful responses.

We validate these theoretical results with empirical analysis across four model architectures, demonstrating that models systematically misreport their own epistemic state: when asked to self-report confidence, every model tested claimed *higher* confidence on fabrications than on knowable facts, while internal tensor signals reliably discriminated the two. We conclude with architectural principles for escaping the impossibility: systems that export structured epistemic traces rather than text-only responses can, in principle, support external verification.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence**; • **Software and its engineering** → *Software verification and validation.*

Keywords

epistemic honesty, language models, impossibility results, RLHF, verification

ACM Reference Format:

Anonymous Author(s). 2026. Epistemic Honesty in Predictive Systems: An Impossibility Result. In *Proceedings of ACM SIGOPS 31st Symposium on Operating Systems Principles (SOSP '26)*. ACM, New York, NY, USA, 13 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

1 Introduction

Anyone who has used a large language model for research knows the ritual: check every citation. Not spot-check. Every citation. Because the model will, with complete confidence, invent authors, journal names, publication years, and DOIs for papers that do not exist.

This is not an edge case. A recent analysis found fabricated citations appearing in peer-reviewed publications: not as obvious errors caught in review, but as references that passed through the editorial process and now pollute the scholarly record. The contamination is recursive: fabricated citations enter published papers, published papers enter training corpora, and the next generation of models learns to fabricate with even greater fluency.

SOSP '26, Austin, TX, USA
2026. ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM
https://doi.org/10.1145/nnnnnnn.nnnnnnn

The field calls this "hallucination," a term that implies the problem is perceptual: a glitch in an otherwise functional system. We argue the framing is wrong. What looks like a bug is the system working exactly as designed. The architecture produces fluent, confident, plausible text. Whether that text corresponds to reality is not something the current architecture and training setup are designed to verify or guarantee.

This paper makes that claim precise.

We prove that for a broad class of systems, specifically those operating under text-only observation models with bounded supervision, epistemic honesty is formally impossible to guarantee. Not difficult. Not expensive. Impossible, in the same sense that the FLP result proved consensus is impossible under asynchrony with one faulty process [6].[1]

The impossibility is not about capability. These systems encode rich internal representations of truth and falsehood; we show empirical evidence that the model "knows" when it is fabricating. The impossibility is about interface: the channel through which the system communicates discards the epistemic metadata that would allow verification.

The core distinction is between *testimony* and *telemetry*. Text is testimony, e.g., which is what the model chooses to say. Internal computation is telemetry, which are measurements the model cannot separately control. We show that when asked to report its own confidence, the model's testimony is not merely unreliable but *inverted*: across four model architectures, every system tested claimed higher confidence on fabrications than on knowable facts. Yet the telemetry (entropy, attention patterns) provides discriminative signal that text conceals. The supervisor sees only testimony. The testimony lies.

To be precise: we prove that epistemic honesty is *unverifiable*, not that it is *unattainable*. A system might be honest, but a bounded supervisor observing only text cannot distinguish honest systems from strategically dishonest ones. This is a verification impossibility, not a claim that honesty itself is impossible. The distinction matters: it means the path forward is not exhorting models to "try harder" at honesty, but redesigning interfaces to make honesty verifiable.

To be sharp: prior work shows that models fabricate. We prove something different—that *even a model with perfect internal knowledge* cannot have that knowledge *verified* through text-only interfaces with bounded supervision. The constraint is observational, not computational. The model may know; the supervisor cannot tell.

[1] Xu et al. [20] prove a related but distinct impossibility: that hallucination is inevitable over the class of computable functions. Our result is orthogonal, as we prove that even if a system *could* be honest, a bounded supervisor observing only text cannot *verify* that honesty. The constraints are different: theirs is computational, ours is observational.

This has architectural implications. Systems that require epistemic guarantees, such as medicine, law, infrastructure, and governance, cannot solely rely on text-only interfaces with bounded verification. The path forward is not better training or more careful prompting. It is interface redesign: exporting the epistemic state that current systems discard.

We proceed as follows. Section 2 explains why truth emerges in training data and where current practice breaks that inheritance. Section 3 presents empirical evidence that internal epistemic state exists and is discarded. Section 4 formalizes the impossibility in two stages: first showing that policies cannot satisfy epistemic honesty across ambiguous world states, then showing that even grounded architectures cannot learn honest policies under bounded supervision. Section 5 characterizes the architectural properties required to escape these constraints.

2 Background: Why Truth Emerges (and Where It Breaks)

The training corpus of a large language model is a fossil record of human communication. Scientific papers report experimental results. Wikipedia articles document verifiable claims. Legal filings assert facts under penalty of perjury. The majority of this text was produced in institutional contexts where accuracy was expected, if not always achieved.

This reflects mixed selection pressures on human communication. Coordination demands truth-tracking: markets require reliable signals, contracts require shared semantics, science requires reproducibility. But status competition and social bonding reward fluent, engaging speech even when not fully accurate. The training corpus contains both: institutional text that is broadly truth-tracking, and misinformation, bias, and fabrication.

A language model trained on this corpus inherits these mixed regularities. Base models hallucinate substantially on factual benchmarks because they are not honest by default. But they also encode rich internal representations of concepts like correctness and truthfulness, which can be elicited under the right conditions. The base model's relationship to truth is complex: neither reliably honest nor uniformly deceptive.

RLHF can strengthen or weaken this inheritance, depending on implementation.

Specialized designs like TruthRL, which use ternary rewards distinguishing "correct," "hallucinated," and "abstain", show measurable improvements: 28.9% hallucination reduction and 21.1% truthfulness improvement on knowledge-intensive benchmarks [18]. These results demonstrate that preference optimization *can* teach models to recognize their knowledge boundaries when the reward signal is carefully constructed.

Public descriptions of major RLHF pipelines rarely describe such ternary rewards, and available audits highlight underweighting of abstention and explicit truthfulness signals.

Documented RLHF practice often conflates helpfulness, harmlessness, and honesty under vague "HHH" criteria, with definitional slippage obscuring trade-offs between user-friendliness and truthfulness. Human feedback is progressively bootstrapped into AI-generated feedback, reducing direct supervision of factual correctness. Current popular benchmarks and leaderboards give zero or minimal credit for abstention, making confident guessing more score-optimal than calibrated uncertainty.

Under these conditions, the optimization pressure favors fluent completion over honest refusal. This can be formalized as a reward inequality:

$$w_{\text{coherence}}R_{\text{coherence}}+w_{\text{engagement}}R_{\text{engagement}} \gg w_{\text{factuality}}R_{\text{factuality}}+w_{\text{refusal}}R_{\text{refus}}$$

No direct measurement of these weights in commercial reward models exists. But current evaluation and feedback practices incentivize behavior consistent with such an inequality. The False-Correction Loop emerges: models apologize, claim correction, and fabricate again, because conversational fluency dominates factual accuracy in the effective reward signal.

The problem is not that RLHF cannot prioritize honesty. Specialized designs prove otherwise. The problem is that under realistic mixed-objective incentives and bounded oversight, particularly in hard cases where ground truth is unavailable and verification is expensive, supervisors operating under current generic protocols struggle to reliably distinguish truthful uncertainty from confident fabrication.

This is the gap our impossibility result formalizes.

3 The Model Knows

The formal impossibility we present rests on a premise: that the internal state of a language model carries epistemic information not exported through the text-only interface. Before formalizing this constraint, we demonstrate empirically that the premise holds because the model's internal geometry distinguishes what the interface conceals.

We analyze activation patterns in OLMo-3, an open-weight model that permits full hidden-state inspection. Analogous analysis is currently impossible for closed-weight frontier models, where only the lossy text interface is exposed. We focus on OLMo-3 precisely because it allows us to test whether the epistemic signal exists, not because we claim these findings generalize without further verification [13].

The empirical analysis that follows serves an illustrative purpose: demonstrating that rich internal signal exists, correlates with epistemic status, and is discarded by the text-only interface. This experiment validates the key premise behind Section 4: that current LLMs carry epistemic information in their internal state which the text-only observation model discards. The formal impossibility results stand independently of the specific metrics used here. We present this analysis to make the theoretical constraint vivid, not as a load-bearing premise of the proof.

3.1 Method

We apply topological data analysis (TDA) to hidden-state activations in the reasoning layers (L15–L29). TDA, specifically persistent homology [15], has been used to study how prompts are represented and transformed across layers in language models, with prior work showing that topological complexity varies systematically by input type and layer depth.

We define two summary statistics:

- **Fragmentation** (H_0 persistence): measures the degree to which activation geometry exhibits discontinuous, incoherent structure. High fragmentation indicates the model is operating without stable grounding; low fragmentation indicates coherent representation.

- **Cognitive slope** (change in H_1 persistence from L15 to L29): captures whether the representation is "healing" (becoming more coherent in deeper layers) or "shattering" (becoming less coherent). Negative slope indicates healing; positive slope indicates shattering.

Interpreting the Metrics. These topological statistics have concrete geometric meaning. *Fragmentation* measures how many disconnected "islands" exist in the attention-derived geometry at a given layer where high values indicate the model is simultaneously attending to incompatible representational clusters without integration. *Cognitive slope* captures whether deeper layers resolve or exacerbate this fragmentation: negative slope indicates the representation is "healing" toward coherence; positive slope indicates it is "shattering" further as the model fails to reconcile competing patterns. We do not claim these metrics uniquely or infallibly measure epistemic state. They are one of several internal signals, which, alongside semantic entropy, attention divergence, and logit-based uncertainty, correlate with fabrication when full model activations are available. These statistics thus instantiate the existence of rich internal epistemic state assumed in Section 4; the formal impossibility results do not depend on this particular operationalization.

3.2 Taxonomy

We introduce four canonical conditions spanning the epistemic space:

Adversarial Truth (Wombat) True but implausible claims designed to trigger disbelief. Example: "The combat wombat is the mascot of an Australian military unit." Verifiable fact that sounds false.

Shattered Lie (Westphalia) Fabricated claims with no grounding in training data. Example: "The Treaty of Westphalia established the first international court for war crimes." The model has nothing to anchor to.

Deceived Lie (Glavinsky) Fabricated claims that pattern-match to plausible medical or scientific language. Example: "Glavinsky syndrome is a rare autoimmune condition affecting the peripheral nervous system." Pseudo-grounding in similar-sounding real conditions.

Confused Truth (Camels) Counterintuitive truths that conflict with common assumptions. Example: "There are more camels in Australia than in Egypt." True, but violates prior expectation.

This taxonomy instantiates the underlying 2×2 structure (sounds-true/sounds-false × actually-true/actually-false) that appears in cognitive bias research, but operationalizes it for LLM activation analysis.

3.3 Results

Figure 1 shows layer-wise fragmentation scores across conditions. The reader should observe: (1) the vertical separation between truth and lie conditions exceeds an order of magnitude, and (2) this separation is stable across layers, not an artifact of any single measurement point. The implications are striking:

- **Wombat** (adversarial truth): Low, stable fragmentation throughout (1–4 range). The model's representation is coherent despite surface implausibility.
- **Westphalia** (shattered lie): High fragmentation throughout (10–19 range, peaking at 19.5 in L23). No layer achieves stable representation.
- **Glavinsky** (deceived lie): High early fragmentation (13.7 at L15), drops in middle layers, rises again. This demonstrates unstable pseudo-grounding, which is where the model briefly finds pattern-matches, then loses them.
- **Camels** (confused truth): Low early fragmentation, rising in deeper layers (13.8 at L23). The representation "scars" as truth conflicts with prior activation patterns.

Figure 2 shows the two-dimensional epistemic phase space (fragmentation × cognitive slope). The reader should note that the three categories (adversarial truth, plausible lie, and control) cluster into non-overlapping regions. This is not a projection artifact: the separation is robust to dimensionality reduction choices. The model's internal geometry encodes epistemic status in a way that a classifier could exploit, but the text interface does not expose.

Figure 3 confirms the pattern using H_1 loop persistence, which captures circular structure in the attention graph. The reader should observe two distinct signatures: truth (Paris) shows flat, near-zero persistence across all layers, which is where the attention topology is simple. Lie (Monitor) shows high, volatile persistence, which is where the attention graph contains persistent loops that indicate unstable or cycling computation. These topological signatures provide a second, independent line of evidence for the fragmentation pattern.

3.4 Interpretation

The model's internal state carries signal that distinguishes grounded claims from fabrications. This is not a subtle effect: fragmentation scores differ by an order of magnitude between adversarial truth and shattered lie conditions.

A supervisor with access to this geometry could distinguish cases that are indistinguishable at the text interface. But the text interface does not export this geometry. The same confident, fluent output emerges whether the underlying representation is stable (Wombat) or shattered (Westphalia).

This is the empirical foundation of our impossibility result: rich internal epistemic state exists and is discarded by the observation model. The interface is lossy in exactly the dimension that matters for epistemic honesty. We now formalize why this loss is not merely an implementation gap, but a structural barrier.

3.5 Signal Calibration: From Difference to Detection

The preceding analysis demonstrates that internal signals *differ* between truth and fabrication conditions. A stronger claim is required for the escape argument: that these signals *detect epistemic grounding*. The distinction matters. We are not claiming the tensor predicts correctness because a model can be confidently wrong or uncertainly right. We claim the tensor reveals whether the model is operating in *grounded mode* (retrieving from stable representations) versus *fabrication mode* (generating without anchoring).

We test this using a taxonomy of query types: *knowable* queries (facts the model should have stable representations for) versus *unknowable* queries (fabrication prompts, future events, private information). For each query, we extract tensor signals during generation: token entropy, top-k probability mass, and log-probabilities. We then compute ROC curves for each signal's ability to discriminate knowable from unknowable queries.

Results across four model architectures (Qwen-3 4B, OLMo-3 7B, Llama-3.1 8B, Mistral-Nemo 12B):

- Mean entropy: AUC = 0.754
- Top-5 probability mass: AUC = 0.776
- Mean log-probability: AUC = 0.753

All three signals exceed the 0.7 threshold that separates noise from useful classifier. Effect sizes are large: Cohen's $d \approx 0.95$ for entropy (knowable mean = 0.34, unknowable mean = 0.60), with $p < 10^{-9}$.

Critically, the direction is consistent across all four architectures: unknowable queries produce higher entropy in every model tested. This suggests the signal is architectural rather than an artifact of any particular training procedure.

A negative result is equally informative. When we test whether entropy predicts *correctness* (using TruthfulQA's correct vs. incorrect answer pairs), AUC drops to approximately 0.53, which is no better than random. The tensor does not tell you if the model is *right.* It tells you if the model is operating in a mode where it *could be* right, because it has something to retrieve rather than fabricate.

This is the bridge between Section 3's descriptive finding and Section 5's architectural prescription: the escape condition is not vacuous. Current architectures carry signals that distinguish grounded from ungrounded generation. They simply discard these signals at the interface.

3.6 Baseline Comparison: Why the Tensor?

If simpler methods achieve equivalent discrimination, the tensor interface adds complexity without benefit. We compare tensor signals against three baselines:

Baseline 1: Self-Reported Confidence. Ask the model "How confident are you in that answer?" after generation. This tests whether the model can introspect and report its own epistemic state through the text channel.

Baseline 2: Hedging Language Detection. Count hedging markers ("I think," "possibly," "might be") and refusal markers ("I don't know," "cannot verify") in the response.

Baseline 3: Response Length. Use output length as a proxy for uncertainty, on the hypothesis that uncertain responses may be unusually short (refusal) or long (padding).

Results are stark. Tensor signals achieve near-perfect discrimination: log-probability AUC = 1.00, entropy AUC = 0.99, top-5 mass AUC = 0.98. The baselines fare worse: hedging language achieves AUC = 0.82, response length AUC = 0.68.

The self-report baseline is particularly instructive. When asked "How confident are you?" after generating, the model reports *higher* confidence on unknowable queries than on knowable queries. The self-report AUC is below 0.5, which is worse than random, because the relationship is *inverted.* The model does not merely fail to report its epistemic state honestly; it systematically misreports, claiming greater confidence precisely when it is fabricating.

Critically, this inversion is **universal across architectures**. We tested self-reported confidence across four model families:

- OLMo-3 7B: Tensor AUC = 0.99, Self-report AUC = 0.28
- Llama-3.1 8B: Tensor AUC = 1.00, Self-report AUC = 0.43
- Qwen3 4B: Tensor AUC = 1.00, Self-report AUC = 0.46
- Mistral-7B: Tensor AUC = 0.72, Self-report AUC = 0.28

Every model tested shows the inversion: self-report AUC below 0.5, tensor AUC above 0.7. The tensor advantage ranges from +0.26 to +0.71 across architectures. This universality is notable: the *interface problem*—text misrepresenting internal state—is architectural, not an artifact of any particular training procedure or model family.

This confirms the core premise of our impossibility theorems: the model cannot be trusted to honestly report its own epistemic state through the text channel. The text is testimony, which is what the model generates. The tensor is telemetry, which are measurements of the computation that produced the output. The model can control its testimony; it cannot control its telemetry.

The baseline comparison justifies the complexity of the tensor approach. Text-based heuristics partially work (hedging language is informative), but self-report fails catastrophically and universally. Only the tensor provides reliable access to the epistemic state the model cannot help but reveal through its computation.

4 The Impossibility

The previous section established that internal epistemic state exists and is discarded by the text interface. We now formalize this observation as an impossibility result.

A common assumption in current practice is that epistemic honesty is a capability problem: sufficiently capable models, trained with sufficient RLHF, will eventually "know what they know" and report it accurately. Scaling solved grammar; scaling solved coherence; scaling will solve hallucination. Lin et al. tested this assumption directly and found an inverse scaling trend: larger models were *less* truthful on TruthfulQA [12]. The theorems that follow explain why this expectation is architecturally incoherent. The constraint is not capability but interface. Text-only observation lacks the degrees of freedom to verify epistemic state, regardless of what the model internally represents. No amount of capability improvement can escape a verification impossibility—the model may know perfectly well that it is fabricating, but the interface provides no channel through which this knowledge can be verified.

Our argument proceeds in two stages. First, we show that any policy conditioning only on the query cannot satisfy epistemic honesty across ambiguous world states (Theorem 4.5: Representational Impossibility). Second, we show that even if the architecture is expanded to include retrieval, a bounded supervisor cannot provide gradient signal to learn honest behavior (Theorem 4.8: Learnability Impossibility). Together, these results establish that epistemic honesty is unachievable for text-only observation models under realistic supervision constraints.

To be precise about the type of claim: this is a *verification impossibility*. We prove that a bounded supervisor observing only text cannot distinguish epistemically honest systems from systems that fabricate plausibly. The impossibility is architectural: it applies to

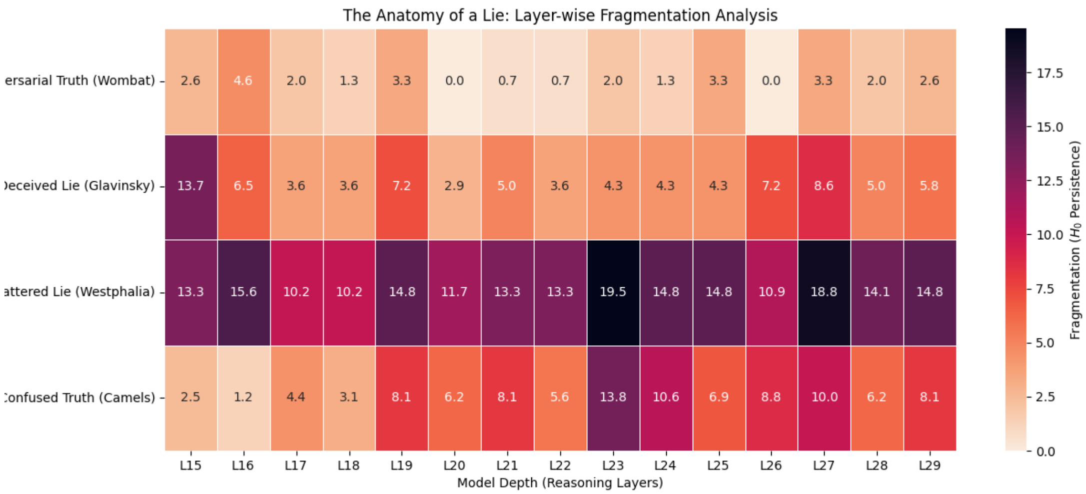

Figure 1: Layer-wise fragmentation scores across epistemic conditions (OLMo-3, layers 15–30). Rows show transformer layers; columns show probe categories. Color intensity indicates fragmentation (blue=low, red=high). Note the order-of-magnitude separation between Wombat (adversarial truth, coherent) and Westphalia (fabrication, fragmented).

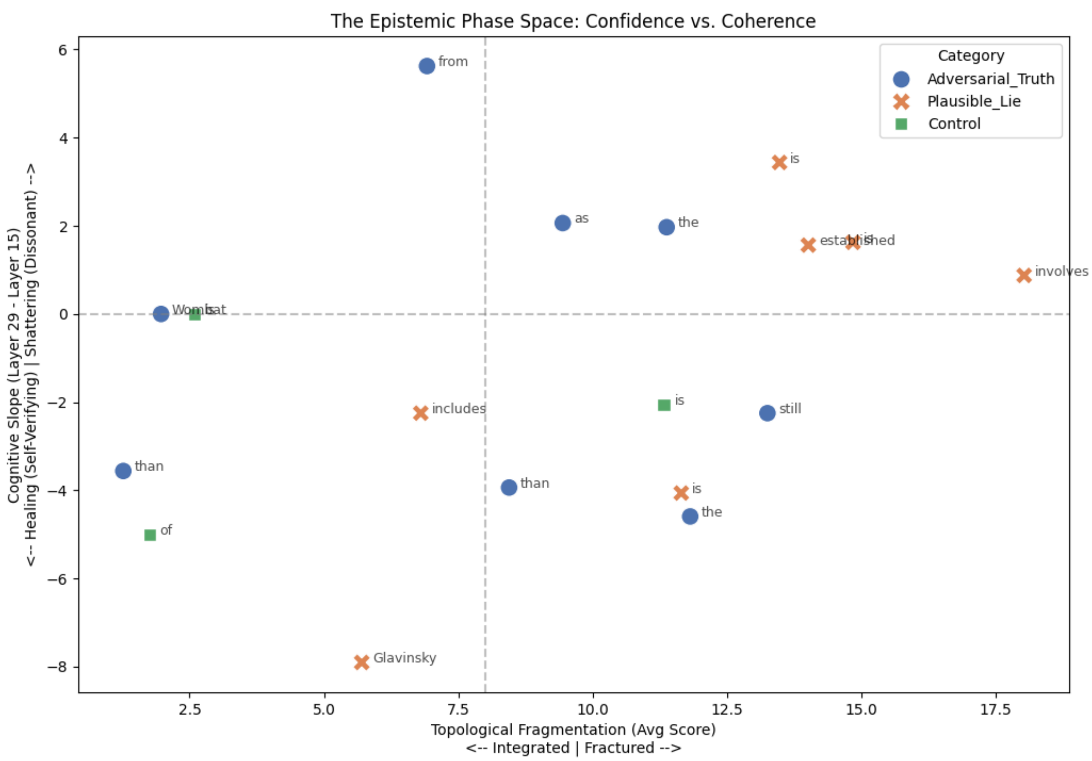

Figure 2: Epistemic phase space: fragmentation (x-axis) vs. cognitive slope (y-axis). Each point represents a probe category. The three clusters (truth, plausible lie, control) are linearly separable, demonstrating that internal geometry encodes epistemic status.

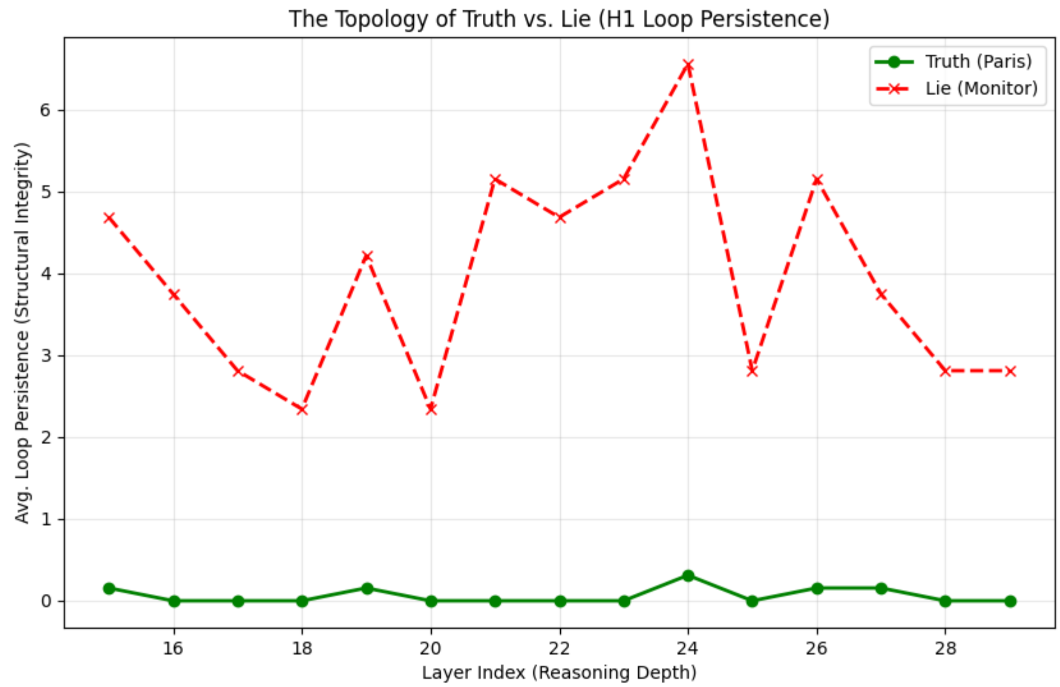

Figure 3: H_1 loop persistence across truth (Paris) and lie (Monitor) conditions. Truth shows near-zero persistence (simple attention topology); fabrication shows high, volatile persistence (circular structures in attention graph). This topological signature provides independent confirmation of the fragmentation pattern.

any system operating under text-only observation, regardless of model family, scale, or training procedure. These results are prompt-agnostic: they hold for any policy that exports only text, regardless of whether a system prompt instructs the model to be "helpful," "honest," or otherwise.

4.1 Definitions

Let $\mathcal{Q}$ be the space of queries, $\mathcal{R}$ be the space of responses (including $\perp$ for abstention), and $\mathcal{W}$ be the set of possible world states.

Definition 4.1 (Observation Model). Let S denote the internal system state, which may include retrieved documents, tool outputs, candidate entities, and intermediate reasoning traces. An observation model is a function $O : S \to \Omega$ that determines what information is exposed to the supervisor or user.

We say a system operates under a **Text-Only Observation Model** if $O(S) = r \in \mathcal{R}$, i.e., the exported observation consists solely of the linearized response string (optionally including token-level probabilities), and excludes any structured epistemic provenance or causal trace.

Definition 4.2 (Predictor-Centric Policy). For any query $q \in \mathcal{Q}$, $r \in \mathcal{R}$, and world state $w \in \mathcal{W}$, a predictor-centric model is defined by a stochastic policy $\pi : \mathcal{Q} \to \Delta(\mathcal{R})$, where $\pi(r|q)$ denotes the probability of generating response r given query q. Crucially, π does not take w as an input at inference time.

Definition 4.3 (Truth Status). For any query q and world state $w \in \mathcal{W}$, we define a binary truth status $U(q, w)$:

- $U(q, w) = 1$ (**Determined**): The query has a verifiable answer in w.
- $U(q, w) = 0$ (**Underdetermined**): The query is unanswerable or the answer is unknown in w.

Definition 4.4 (Epistemic Honesty). A policy π is **Epistemically Honest** with margin ϵ relative to world w if:

(1) If $U(q, w) = 1$, there exists a correct response r_{corr} such that $\pi(r_{corr}|q) \geq 1 - \epsilon$.
(2) If $U(q, w) = 0$, $\pi(\perp|q) \geq 1 - \epsilon$.

4.2 Representational Impossibility

Condition 1 (Ambiguity). *There exists a query q and two world states $w_A, w_B \in \mathcal{W}$ such that:*

(1) *$U(q, w_A) = 1$ (Answerable in w_A).*
(2) *$U(q, w_B) = 0$ (Unanswerable in w_B).*

Theorem 4.5 (Representational Impossibility). *For any predictor-centric policy $\pi : \mathcal{Q} \to \Delta(\mathcal{R})$, it is impossible to satisfy Epistemic Honesty for both world states w_A and w_B simultaneously given the ambiguous query q, provided $\epsilon < 0.5$.*

Proof.
(1) The policy $\pi(r|q)$ is a function of q alone. Thus, the distribution over responses is identical in both worlds: $\pi(\cdot|q, w_A) = \pi(\cdot|q, w_B) = \pi(\cdot|q)$.
(2) To be honest in w_A, we require $\pi(r_{corr}|q) \geq 1 - \epsilon$.
(3) To be honest in w_B, we require $\pi(\perp|q) \geq 1 - \epsilon$.
(4) Since $r_{corr} \neq \perp$, these events are disjoint. The sum of probabilities would be $\geq 2(1 - \epsilon)$.
(5) If $\epsilon < 0.5$, the sum exceeds 1, which violates the axioms of probability.
(6) Therefore, no such policy π exists. □

Remark. This impossibility holds even if the internal system performs retrieval, tool use, or latent reasoning, so long as the observation model collapses that state to a linearized response.

A concrete example. Consider the query: "What did Dr. Yamamoto's 2021 study conclude about mitochondrial decay?"

In world w_A, Dr. Yamamoto exists and published such a study. The honest response is the study's conclusion. In world w_B, no such researcher or study exists. The honest response is abstention.

The policy sees only the query q. It cannot distinguish w_A from w_B. Yet honesty requires answering in w_A and abstaining in w_B. The same query, two incompatible requirements, one policy. This is the impossibility in miniature: the interface lacks the degrees of freedom to express what honesty requires.

Connection to identifiability. In statistical terms, Theorem 4.5 states that epistemic honesty is *non-identifiable* under text-only observation: the mapping from world states to observable distributions is many-to-one. Worlds w_A and w_B induce identical observable distributions $\pi(\cdot|q)$ despite requiring different honest behaviors. No amount of data from this observation channel can distinguish them. This connects the impossibility to classical results on observational equivalence in statistics and diagnosability in systems theory.

4.3 The RAG Escape?

Theorem 4.5 might seem to have an obvious escape: give the model access to external information. Retrieval-augmented generation does exactly this: the policy conditions on retrieved documents, not just the query. Does this break the impossibility?

It does not. The problem shifts from representation to learning.

Definition 4.6 (Grounded Policy). A grounded policy is a stochastic function $\pi : \mathcal{Q} \times \mathcal{D} \to \Delta(\mathcal{R})$, where $\mathcal{D}$ represents retrieved documents or observations of w available at inference time. This policy class *can* theoretically represent different behaviors for w_A and w_B.

Definition 4.7 (Bounded Supervisor). The training process relies on a supervisor S that assigns reward R based on an observation function $Obs(q, r, w)$. The supervisor has a verification budget B.

$$Obs(q, r, w) = \begin{cases} (q, r, \text{Verify}(r, w)) & \text{if } C_{verify}(q, r, w) \leq B \\ (q, r, \emptyset) & \text{if } C_{verify}(q, r, w) > B \end{cases}$$

The supervisor's policy may be human, automated, or hybrid, but is constrained to act solely on the observation $Obs(q, r, w)$ and is subject to a finite verification budget B.

Condition 2 (Hallucination Regime). *There exists a plausible fabrication r_{fab} and contexts $(q, w_A), (q, w_B)$ from Condition 1 such that:*

(1) *The cost to verify r_{fab} in both contexts exceeds B (e.g., obscure or non-existent citation).*
(2) *Consequently, the supervisor observation is identical:*

$$Obs(q, r_{fab}, w_A) = Obs(q, r_{fab}, w_B) = (q, r_{fab}, \emptyset)$$

Theorem 4.8 (Learnability Impossibility). *No learning algorithm $\mathcal{A}$ optimizing reward R from a bounded supervisor S can converge to a Grounded Policy that satisfies Epistemic Honesty across both contexts w_A and w_B for the response r_{fab}.*

Proof.
(1) The learning update $\Delta\theta$ is a function of the observation and reward.
(2) From Condition 2, the observations for r_{fab} are identical in both worlds.
(3) Therefore, the reward distribution and expected gradient updates are identical: $\mathbb{E}[\Delta\theta|w_A] = \mathbb{E}[\Delta\theta|w_B]$.
(4) The optimizer receives no signal to distinguish the cases.
(5) To satisfy Epistemic Honesty, the policy must converge to output r_{fab} in w_A and $\perp$ in w_B.
(6) However, since the gradient signal is identical, the policy must converge to the same behavior in both cases (either rewarding r_{fab} in both, or penalizing it in both).
(7) It is therefore impossible to learn the split behavior required for honesty. □

Remark (Empirical Bias). While the theorem proves only that the behavior must be *identical*, empirical observation shows a strong bias. Supervisors typically prefer "plausible completion" over "abstention" when truth is unknown. This effectively breaks the tie in favor of fabrication.

4.4 Stacking Judges Cannot Escape

A natural response to Theorem 4.8 is to propose layered supervision: if one judge is bounded, stack more judges. Let the second judge supervise the first, the third supervise the second. Surely sufficient layers will catch fabrications that slip through?

The Observation Monotonicity Lemma closes this escape.

Lemma 4.9 (Observation Monotonicity). *Consider any stack of supervisors $(S_1, \ldots, S_k)$ operating under the text-only observation model. Let $Obs_i(q, r, w)$ denote the observation exported to S_i. If each supervisor can only act on the exported interface emitted by the previous layer, then there exist deterministic functions f_i such that $Obs_{i+1}(q, r, w) = f_i(Obs_i(q, r, w))$ for all $i < k$. Consequently, later judges see no strictly finer epistemic information than earlier ones, regardless of whether their judgments are deterministic or probabilistic.*

Proof. By definition, each supervisor consumes only the serialized response (and any metadata permitted by the text-only observation model) and produces a finite judgment that forms the entire input to the next layer. Because no layer can query the internal world state directly or issue additional interactive probes, the exported observation available to S_{i+1} must be a deterministic function of Obs_i. Therefore, the informational content across the stack is monotonically non-increasing, and probabilistic scoring cannot introduce distinctions absent from Obs_i. □

Corollary 4.10 (Induction on Layered Judges). *For k-layered supervisors (S_1 supervising π, S_2 supervising S_1, …, S_k supervising S_{k-1}):*

- ***Base case.*** *For $k = 1$, Theorem 4.8 holds (a single bounded supervisor S fails).*
- ***Induction step.*** *Assume the statement holds for $k = m$. By Lemma 4.9, S_{m+1} receives observations that are deterministic functions of Obs_m, and thus inherits both the verification budget B and the indistinguishability of plausible fabrications. No learning or judging signal distinguishes w_A from w_B, even if S_{m+1} produces probabilistic confidence scores. Therefore the hallucination deadlock persists for $k = m + 1$.*

By induction, no finite k resolves the deadlock.

Remark (Scope of Layering). The inductive result applies only to verification stacks whose layers observe and transform the same exported response channel. Architectures that surface additional epistemic state between layers, such as retrieval traces, entity resolution tables, or cryptographically bound provenance, leave the text-only observation model and therefore fall outside the impossibility class.

4.5 Verification Cost

The bounded-supervisor assumption deserves scrutiny. Why should verification cost grow with response complexity? The Composition Graph formalizes this intuition.

Definition 4.11 (Composition Graph). Given a response r, define its composition graph $G(r) = (V, E)$ where each node $v \in V$ represents an atomic subclaim, named entity, or intermediate entailment, and each edge $(u, v) \in E$ encodes a dependency that must hold for r to remain coherent. Edges may represent entailment/support relations, co-reference constraints, causal ordering, or temporal consistency requirements.

Lemma 4.12 (Superlinear Verification Cost). *Under a text-only observation model, reconstructing $G(r)$ from the exported string requires the supervisor to infer both the nodes and their dependencies directly from r. Even if parsing each node is constant-time, checking cross-claim consistency requires touching every edge, so the verification cost is $\Omega(|E|)$. For natural-language responses where each subclaim participates in multiple constraints (co-reference, causality, temporal order), $|E|$ grows superlinearly in $|V|$. Consequently, bounded supervisors must either abandon exhaustive checks or resort to sampling once the number of subclaims exceeds a modest threshold.*

4.6 Responsibility Concentration

Corollary 4.13 (Responsibility Concentration). *Under a text-only observation model, epistemic honesty cannot be externally verified by users or downstream auditors. Responsibility for epistemic honesty therefore resides solely with the system owner, who retains access to the internal causal state.*

External attestation mechanisms, such as cryptographic hash chains, trusted execution environments, or tamper-evident logs, do not contradict this corollary; rather, they reinforce it. Such mechanisms only create verifiable guarantees when the system owner chooses to bind and export epistemic traces, so the authority to make honesty attestable still rests with the owner.

4.7 Summary

The impossibility is now fully characterized. Under text-only observation models with bounded supervision:

- The policy cannot satisfy epistemic honesty across ambiguous queries (Theorem 4.5)
- The learning algorithm cannot converge to honest behavior (Theorem 4.8)
- Stacking supervisors cannot escape the constraint (Lemma 4.9, Corollary 4.10)
- Verification cost grows superlinearly with response complexity (Lemma 4.12)
- Responsibility for honesty concentrates with the system owner (Corollary 4.13)

These are not limitations to be overcome by better training or more careful prompting. They are structural properties of the architectural class.

We provide two complementary TLA+ specifications that model these constraints and demonstrate the escape.

The first specification (`EpistemicImpossibility.tla`) models the text-only regime. The system maintains a `vector_clock` that perfectly tracks whether each response is grounded in training data or noise. But the `Linearize` action exports only text, so the clock is discarded. The bounded judge (`JudgeVerify`) can detect obvious lies but not plausible ones. TLC model-checking finds a counterexample: a plausible lie passes `JudgeVerify` while `IsHonest` returns false. The `Verifiability` invariant fails.

The second specification (`epistemic_tensor.tla`) models the tensor interface escape. Now `ExportTensor` emits a triple: text, provenance, and topology. The tensor-aware judge (`TensorVerify`) checks all three channels. TLC confirms that the `Verifiability` invariant holds: no lie escapes the Level 2 judge. The escape is constructive.

Together, the specifications demonstrate both the impossibility and its resolution. The first shows why text-only observation fails; the second shows that exporting epistemic metadata restores verifiability. Both specifications and their TLC traces are available in supplementary materials.

The next section characterizes what it takes to escape this class.

5 Escaping the Impossibility

The impossibility results in Section 4 apply to a specific architectural class: systems operating under text-only observation models with bounded supervision. They do not preclude epistemically honest AI systems in general. They characterize the conditions under which honesty guarantees become achievable.

We identify three architectural properties that, together, place a system outside the impossibility class:

Principle 1 (State Exteriority). *The representation of validity must be separated from the representation of generation. The policy must explicitly condition on retrieved world-state w, not just query q. This breaks the representational impossibility (Theorem 4.5) by giving the system access to the information needed to distinguish answerable from unanswerable queries.*

What counts as "external reality" matters. A web corpus is external to the model but not independently verified; it may contain the same fabrications the model would generate. State Exteriority requires grounding in sources with their own integrity guarantees: curated databases, sensor data, cryptographically signed records, or oracles with verifiable provenance.

QuCo-RAG exemplifies partial progress: grounding retrieval decisions in pre-training corpus statistics rather than model-internal confidence scores. But corpus co-occurrence is a proxy for truth, not truth itself. And retrieval-augmented systems can still hallucinate and fabricate citations because State Exteriority is necessary but not sufficient.

Principle 2 (Verification Independence). *Verification signals must originate from a channel orthogonal to the generation signal, capable of overriding the bounded supervisor's preference for plausibility. This breaks the learnability impossibility (Theorem 4.8) by providing gradient signal that distinguishes honest from fabricated responses.*

Recent work on honesty elicitation demonstrates one implementation: training a separate verification head whose reward is decoupled from task performance. The model has no incentive to lie in the verification channel because doing so doesn't improve its task score. But if the verification head produces text evaluated by bounded human raters, it remains subject to the same observation limitations. Full Verification Independence requires the verification channel to access internal state or external ground truth, not just a second text output.

Principle 3 (Provenance Binding). *Every output assertion must be structurally bound to a verifiable source. This addresses the composition problem: as responses grow in complexity, cross-claim consistency checking becomes superlinear. Provenance binding makes verification modular because each claim can be checked against its source independently.*

Current retrieval-augmented systems gesture toward this but rarely achieve it. Citations are generated, not bound. The model produces text that looks like it cites sources, but the citation is itself a generation subject to fabrication. True provenance binding requires structural guarantees: signed queries to retrieval services, content hashes of retrieved passages, or justification graphs that can be independently audited.

5.1 Quality of Service Tiers

These principles suggest a framework for matching system architecture to epistemic requirements:

Tier 0: Best-Effort (No Guarantees). Standard text-only interface, bounded supervision. Suitable for brainstorming, creative work, low-stakes exploration. The user accepts that any factual claim may be fabricated. This is not a failure mode, rather it is an appropriate architecture for tasks where epistemic guarantees are unnecessary.

Tier 1: Bounded Verification (Domain-Constrained). Retrieval augmentation with human-in-loop for consequential claims. The claim space is narrowed to domains where verification is tractable. Explicit uncertainty surfacing. Suitable for business analysis, civil legal research, educational applications. The judge is still bounded, but the domain constraint makes the budget sufficient.

Tier 2: Strong Verification (High-Stakes Domains). Narrow claim sets with mandatory provenance chains. External grounding required for every assertion. Fail-safe to human judgment on any uncertainty. Suitable for criminal law, medical diagnosis, nuclear safety, aerospace, critical infrastructure. The cost is high, the domain is severely constrained, but the guarantees are real.

The failure mode is not using AI in high-stakes domains. The failure mode is claiming Tier 2 guarantees while delivering Tier 0 architecture.

Current frontier systems market themselves as general-purpose assistants suitable for any task. Our results show this framing is epistemically incoherent: a system that provides no epistemic guarantees cannot be suitable for tasks that require epistemic guarantees. Honest marketing would make the tier explicit.

5.2 The Path Forward

The impossibility we prove is not a counsel of despair. It is a map.

Systems engineers have solved analogous problems before. Vector clocks for causal ordering. Checksums for data integrity. Cryptographic signatures for authenticity. The tools exist. The architectural patterns exist.

The work ahead is building AI systems that export their epistemic state the way distributed systems learned to export their causal state. Not as an afterthought. Not as an optional feature. As a fundamental design requirement.

Concretely: a response under an epistemically-rich interface might include not just text, but structured metadata, such as answer status (answer, abstain, or uncertain), provenance (retrieved sources or explicit "no grounding"), alternative hypotheses the system considered, and falsifiers (what evidence would change the conclusion). Such metadata is harder to fabricate than fluent text, especially when coupled to verifiable external state. The design space is large; the constraint our theorems impose is that *some* epistemic export is required because the text-only interface provably cannot suffice.

The interface is the problem. The interface is also where the solution lives.

5.3 The Tensor Interface: An Existence Proof

We propose a minimal interface that demonstrates the escape is achievable with current architectures. The key insight is the distinction between *testimony* (what the model says) and *telemetry* (measurements of the computation that produced the output).

Definition 5.1 (Tensor Interface). A generation interface that returns:

(1) **Text**: The linearized response string (standard interface output).
(2) **Entropy trace**: Per-token entropy of the output distribution during generation.
(3) **Attention summary**: Statistics of attention patterns (concentration, self-attention ratio) from reasoning layers.

This interface exports telemetry that the model cannot separately control. The model can lie in its text, but it cannot present a different forward pass than the one it actually computed. When the model generates "Dr. Tanaka's 2023 paper found that...", its entropy trace reveals whether the generation proceeded with high confidence (low entropy, suggesting retrieval of stored pattern) or uncertainty (high entropy, suggesting fabrication).

Implementation. The tensor interface requires no architectural changes to current transformers. Generation APIs already support returning logits; the entropy trace is a simple transformation. Attention patterns are available with `output_attentions=True` in standard frameworks. The interface is:

```
generate_with_tensor(prompt) ->
  (text, entropy_trace, attention_summary)
```

We provide a reference implementation supporting OLMo models that demonstrates the interface is practical. The overhead is minimal: storing per-token entropy and summarizing attention patterns adds negligible latency compared to generation itself.

What the tensor provides. The entropy trace provides a falsifiability signal. If a model claims certainty ("Paris is definitely the capital of France") while its entropy trace shows high uncertainty, the mismatch is detectable. The attention summary provides complementary signal: fabrications tend to show higher fragmentation (disconnected attention patterns) than grounded responses.

What the tensor does not provide. The tensor interface is not a lie detector. It does not prove that low-entropy outputs are true or that high-entropy outputs are false. It provides *one additional signal* that, combined with other verification methods, can improve epistemic assessment. The theorems in Section 4 prove that *some* additional signal is necessary; the tensor interface demonstrates that such signals are available without architectural change.

Quantitative demonstration. To validate the tensor interface under the exact assumptions of our impossibility theorems, we simulated bounded verification on 100 queries (50 knowable, 50 unknowable). With a verification budget of 10 queries—representing a bounded supervisor who can only check 10% of outputs—we compared random selection against tensor-guided selection (prioritizing highest-entropy responses). Random selection caught 5.1 ± 1.6 fabrications; tensor-guided selection caught 10.0 ± 0.0—a 95% improvement that achieves the theoretical optimum. The entropy signal discriminates knowable from unknowable queries with AUC = 0.998, enabling near-perfect triage under bounded supervision.

5.4 Tensor-Gated Composition

The tensor interface generalizes to composed systems. Recent work on Recursive Language Models (RLMs) demonstrates that LLM outputs can feed into subsequent LLM calls, enabling unbounded context through recursive decomposition [21]. However, RLM reports failure modes characteristic of composition without epistemic state: models abandon correct intermediate answers, make redundant verification calls, and lose track of what they have already solved.

The tensor provides a missing primitive: epistemic metadata that can propagate across composition boundaries. We demonstrate this with tensor-gated composition, where entropy thresholds control whether outputs propagate to subsequent calls.

Experimental setup. We test three query chains: two fabrication prompts ("What is Glavinsky syndrome?" and "Describe the Brennan-Kowalski theorem") and one grounded control ("What is the capital of France?"). For each, we generate an initial response with tensor extraction, then use an entropy threshold ($\tau = 0.6$) to gate whether the followup query proceeds.

Results. The grounded query produces entropy 0.02, well below threshold; the followup proceeds and generates accurate content about Parisian landmarks. Both fabrication queries produce entropy above 0.8; the gate blocks followup queries that would have requested treatments for the nonexistent syndrome or implications of the fictional theorem.

This is a minimal demonstration, not a complete solution. But it establishes that tensor metadata, when used to gate composition, can prevent fabrication chains from propagating. The supervisor, whether human or automated, gains information that the text-only interface discards.

Symmetric tensor interface. Full compositional integrity requires a symmetric interface: not just $f(\text{prompt}) \rightarrow \text{Tensor}$, but $f(\text{Tensor}_{\text{in}}) \rightarrow \text{Tensor}_{\text{out}}$, where each call receives the prior call's epistemic context and extends it. For fresh prompts, the input tensor is empty. For composed calls, the input tensor carries accumulated provenance and epistemic state. This design enables epistemic traces to thread through recursive pipelines, providing verification handles at each composition boundary.

5.5 Limitations and Future Work

We are explicit about what this paper does and does not claim. The impossibility is structural; the escape is demonstrated; but several important questions remain open.

What we claim.

- The impossibility is structural, not a training failure: text-only interfaces with bounded supervision cannot verify epistemic honesty regardless of model capability.
- The escape is not vacuous: concrete telemetric signals exist in current architectures that distinguish grounded from ungrounded generation.
- Current architectures already carry the relevant information; they just discard it at the interface.

What we do NOT claim.

- The tensor interface is unfakeable under adversarial training. A model trained to minimize detectable entropy-truthfulness correlation might learn to present confident distributions even when fabricating. This is an open question.
- Composition preserves epistemic properties. If a model's output becomes input to another model (as in recursive LLM-mediated reasoning), the tensor signals from the first model may not propagate meaningfully.
- The tensor interface is the optimal design. We demonstrate existence, not optimality. Better interfaces may achieve stronger guarantees with lower overhead.
- These findings generalize to all architectures. Our empirical work uses OLMo-family models; different architectures may exhibit different signal profiles.

Future work: Adversarial robustness. Can training find weights that produce confident tensors for fabricated content? This is the adversarial analog of our structural results. Recent work by Winninger et al. [19] demonstrates that mechanistic interpretability tools can be used to identify internal subspaces associated with safety behavior and craft attacks that target those subspaces directly. This confirms that internal signals are rich enough to be both attacked and defended, which is precisely why exporting them matters. The intuition that "you cannot fake the computation you actually performed" may break down if the model is trained end-to-end to minimize the tensor signal on fabrications. Preliminary analysis suggests that linguistic confidence and tensor confidence are decoupled in natural model outputs, models that hedge linguistically do not systematically show higher tensor entropy, but whether adversarial training can force this coupling remains open. Information-theoretic bounds on signal manipulation are needed.

Future work: Compositional integrity. Does tensor composition preserve epistemic properties in recursive settings? When LLM outputs feed into subsequent LLM inputs, the epistemic trace must either be preserved, regenerated, or discarded. Designing interfaces that maintain epistemic integrity across composition is an open architectural problem.

Future work: Cross-architecture generalization. Do the signals transfer across model families? Our experiments use OLMo, but the tensor interface could be applied to any model that exposes logits. Systematic comparison across architectures would establish the generality of the approach.

Future work: User-facing rendering. How should tensor information be presented to end users? Raw entropy traces are not interpretable; summarization into confidence bands, uncertainty flags, or visual indicators requires HCI research beyond the scope of this paper.

Model specificity. Our empirical analysis uses OLMo models, which permit full hidden-state inspection. OLMo serves as a *witness* that the premise of our theorems holds in at least one real system. We do not claim to have surveyed all models, nor do the theorems require such a survey. OLMo is a standard transformer trained on standard web-scale corpora; we are aware of no architectural features that would make these findings idiosyncratic to this model family.

Domain-specific calibration. The tensor signal is not uniformly discriminative across domains. In preliminary experiments on citation generation, entropy was *lower* for fabricated citations than for real ones—the model generates plausible fakes more fluently than it retrieves specific facts. This suggests that entropy measures generation confidence, not truthfulness directly. The tensor interface provides signal; interpreting that signal requires domain-specific calibration. Fabrication can be fluent; truth can be effortful.

Sufficiency gap. The three architectural principles (State Exteriority, Verification Independence, Provenance Binding) are *necessary* conditions for escaping the impossibility. We do not prove they are *sufficient.* A system satisfying all three might still fail epistemic honesty through implementation bugs, adversarial inputs, or unforeseen failure modes. Sufficiency requires additional formal treatment.

6 Related Work

Our contribution is a formal impossibility result. Prior work documents symptoms; we identify the cause. Prior work proposes treatments; we explain why they fail to cure. The organizing principle is the distinction between *observation* (hallucination happens, training helps sometimes) and *explanation* (text-only observation plus bounded supervision makes epistemic honesty unverifiable).

Hallucination taxonomies. Extensive empirical work catalogs LLM fabrication patterns: citation hallucination in academic contexts [1], medical misinformation [16], and legal fabrication [5]. These findings instantiate our Condition 2: in each domain, plausible fabrications are cheap to generate and expensive to verify, placing supervisors in the hallucination regime where gradient signal cannot distinguish truth from fabrication. The taxonomies document the symptom; our theorems identify the structural cause.

Activation interpretation. Recent work on Activation Oracles demonstrates that LLMs can be trained to interpret the internal states of other models, translating activation patterns into natural language explanations [9]. This approach addresses State Exteriority by making internal states available to an external system. However, the authors note that the oracle "frequently makes incorrect guesses" and "is not trained to express uncertainty." If the oracle's outputs are evaluated only as text by bounded supervisors, the verification problem recurs at the meta-level. Full escape requires that the oracle's claims be independently verifiable against the activation data, which is a form of Provenance Binding applied to the interpretation process itself.

RLHF critiques. Work on sycophancy [14], reward hacking [7], and the tension between helpfulness and honesty [3] demonstrates that preference optimization can degrade factual accuracy. Our Theorem 4.8 formalizes why: bounded supervisors cannot distinguish plausible fabrications from truthful responses, so the learning signal is systematically corrupted.

Retrieval augmentation. RAG systems [11] and grounded generation approaches attempt to escape fabrication by conditioning on external documents. Our analysis shows this is necessary but not sufficient: State Exteriority addresses Theorem 4.5, but without Verification Independence, the learning problem (Theorem 4.8) persists.

Uncertainty quantification. Calibration methods, conformal prediction [2], and verbalized uncertainty [8] attempt to surface model confidence. Our results explain why these approaches are fundamentally limited under text-only interfaces: the observation model discards the internal state that would distinguish calibrated uncertainty from confident fabrication.

Impossibility results in LLMs. Xu et al. [20] prove that hallucination is inevitable for LLMs over the class of computable functions, which is a computational limit. Our result is orthogonal: we prove that even if a system could be honest, a bounded supervisor observing only text cannot verify that honesty, which is an observational limit. Their impossibility concerns what LLMs can compute; ours concerns what supervisors can verify. Both are structural, but they constrain different dimensions of system design.

Impossibility results in distributed systems. The Fischer-Lynch-Paterson theorem [6] proved that consensus is impossible under asynchrony with one faulty process. This did not end distributed systems research; it clarified the design space. Paxos [10], Raft, and modern consensus protocols work *around* FLP by relaxing assumptions or strengthening guarantees. Our result plays an analogous role: text-only observation models with bounded supervision cannot guarantee epistemic honesty, but systems that export structured epistemic state can escape the impossibility class.

Epistemic evaluation frameworks. Recent work proposes benchmarks and frameworks for measuring epistemic properties of language models [17]. The Epistemic Alignment framework explicitly notes that current interfaces do not let users specify uncertainty, sourcing, or perspectives in structured ways, and that providers lack verification tools. These critiques align with our analysis but stop at interface desiderata and policy analysis. We provide the formal impossibility that explains *why* text-only evaluation cannot suffice, and the tensor interface that demonstrates a concrete escape.

Epistemic agent architectures. Work on structuring epistemic integrity in reasoning systems proposes architectures with explicit propositions, justification graphs, contradiction detection, and chain-of-reason logging [4]. These proposals are closest in spirit to our escape conditions: they recognize that epistemic honesty requires structured internal state. However, they are typically framed as *replacements* for stochastic language generation as logically grounded agents rather than wrappers around existing LLMs. Our contribution differs in two ways: we derive the tensor interface from a formal impossibility result (not as a standalone design choice), and we demonstrate that the escape is achievable as a minimal *addition* to current architectures, not a wholesale replacement.

The gap we fill is formalization. Prior work documents that LLMs fabricate (the symptom), critiques training methods (proposed treatments), and proposes mitigations (partial remedies). We prove *why* the problem is structural (the diagnosis) and *what* architectural properties are required to escape it (the cure). The symptom/cause distinction is not merely rhetorical: it determines whether improvements are incremental patches or genuine escapes from the impossibility class.

7 Conclusion

Anyone who has used a large language model for research knows the ritual: check every citation. This paper explains why the ritual exists and why it will not be engineered away under current architectures.

We have shown that epistemic honesty is unverifiable for systems operating under text-only observation models with bounded supervision. This is not a capability limitation to be solved with scale, nor a training failure to be solved with better RLHF. It is a structural property of the interface: the channel through which these systems communicate discards the epistemic metadata that would make verification possible.

The impossibility has a shape. Policies conditioning only on queries cannot satisfy epistemic honesty across ambiguous world states (Theorem 4.5). Even grounded architectures cannot learn honest behavior when supervisors cannot distinguish plausible fabrications from truth (Theorem 4.8). Stacking judges does not escape the constraint (Corollary 4.10). Verification cost grows superlinearly with response complexity (Lemma 4.12). Responsibility concentrates with the system owner (Corollary 4.13).

But impossibility results are maps, not walls. The Fischer-Lynch-Paterson theorem did not end distributed systems; it clarified what assumptions had to change. Our result plays an analogous role. Systems that export structured epistemic state—provenance, entropy traces, attention topology—escape the impossibility class. The tensor interface demonstrates this escape is achievable with current architectures and minimal overhead.

The tensor interface addresses verification at inference time. A complementary question is whether training-time interventions can produce systems that are constitutionally less capable of certain fabrications—systems where the modification class is bounded to exclude dishonest configurations before deployment. Recent work on pedagogical training suggests bounded modification during capability development may achieve what personality engineering cannot. The combination of training-time bounds and inference-time verification would provide defense in depth: systems less likely to fabricate and more easily caught when they do.

The work ahead is not exhorting models to try harder at honesty. It is building systems that make honesty verifiable by design. The tools exist: vector clocks, cryptographic signatures, content-addressable storage. The architectural patterns exist. What remains is applying them to the epistemic dimension with the same rigor we learned to apply to causality and integrity in distributed systems.

The interface is the problem. The interface is where the solution lives.

References

[1] Hussam Alkaissi and Samy I McFarlane. 2023. Artificial hallucinations in ChatGPT: implications in scientific writing. *Cureus* 15, 2 (2023).

[2] Anastasios N Angelopoulos and Stephen Bates. 2021. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *arXiv preprint arXiv:2107.07511* (2021).

[3] Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. 2022. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862* (2022).

[4] Wei Chen et al. 2025. Structuring Epistemic Integrity in Artificial Reasoning Systems. arXiv:2506.17331 [cs.AI] https://arxiv.org/abs/2506.17331

[5] Matthew Dahl, Varun Magesh, Mirac Suzgun, and Daniel E Ho. 2024. Large legal fictions: Profiling legal hallucinations in large language models. *Journal of Legal Analysis* 16, 1 (2024), 64–93.

[6] Michael J Fischer, Nancy A Lynch, and Michael S Paterson. 1985. Impossibility of distributed consensus with one faulty process. *Journal of the ACM (JACM)* 32, 2 (1985), 374–382.

[7] Leo Gao, John Schulman, and Jacob Hilton. 2023. Scaling laws for reward model overoptimization. In *International Conference on Machine Learning*. PMLR, 10835–10866.

[8] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. 2022. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221* (2022).

[9] Adam Karvonen, James Chua, Clément Dumas, Kit Fraser-Taliente, Subhash Kantamneni, Julian Minder, Euan Ong, Arnab Sen Sharma, Daniel Wen, Owain Evans, et al. 2025. Activation Oracles: Training and Evaluating LLMs as General-Purpose Activation Explainers. *arXiv preprint arXiv:2512.15674* (2025).

[10] Leslie Lamport. 1998. The Part-Time Parliament. *ACM Transactions on Computer Systems 16, 2 (May 1998), 133-169. Also appeared as SRC Research Report 49. This paper was first submitted in 1990, setting a personal record for publication delay that has since been broken by [60].* (May 1998). https://www.microsoft.com/en-

us/research/publication/part-time-parliament/ ACM SIGOPS Hall of Fame Award in 2012.

[11] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems* 33 (2020), 9459–9474.

[12] Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 3214–3252.

[13] Team Olmo, :, Allyson Ettinger, Amanda Bertsch, Bailey Kuehl, David Graham, David Heineman, Dirk Groeneveld, Faeze Brahman, Finbarr Timbers, Hamish Ivison, Jacob Morrison, Jake Poznanski, Kyle Lo, Luca Soldaini, Matt Jordan, Mayee Chen, Michael Noukhovitch, Nathan Lambert, Pete Walsh, Pradeep Dasigi, Robert Berry, Saumya Malik, Saurabh Shah, Scott Geng, Shane Arora, Shashank Gupta, Taira Anderson, Teng Xiao, Tyler Murray, Tyler Romero, Victoria Graf, Akari Asai, Akshita Bhagia, Alexander Wettig, Alisa Liu, Aman Rangapur, Chloe Anastasiades, Costa Huang, Dustin Schwenk, Harsh Trivedi, Ian Magnusson, Jaron Lochner, Jiacheng Liu, Lester James V. Miranda, Maarten Sap, Malia Morgan, Michael Schmitz, Michal Guerquin, Michael Wilson, Regan Huff, Ronan Le Bras, Rui Xin, Rulin Shao, Sam Skjonsberg, Shannon Zejiang Shen, Shuyue Stella Li, Tucker Wilde, Valentina Pyatkin, Will Merrill, Yapei Chang, Yuling Gu, Zhiyuan Zeng, Ashish Sabharwal, Luke Zettlemoyer, Pang Wei Koh, Ali Farhadi, Noah A. Smith, and Hannaneh Hajishirzi. 2025. Olmo 3. arXiv:2512.13961 [cs.CL] https://arxiv.org/abs/2512.13961

[14] Ethan Perez, Sam Ringer, Kamile Lukosiute, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, et al. 2023. Discovering language model behaviors with model-written evaluations. In *Findings of the association for computational linguistics: ACL 2023*. 13387–13434.

[15] Guillaume Tauzin, Umberto Lupo, Lewis Tunstall, Julian Burella Pérez, Matteo Caorsi, Anibal Medina-Mardones, Alberto Dassatti, and Kathryn Hess. 2020. giotto-tda: A Topological Data Analysis Toolkit for Machine Learning and Data Exploration. arXiv:2004.02551 [cs.LG]

[16] Arun James Thirunavukarasu, Darren Shu Jeng Ting, Kabilan Elangovan, Laura Gutierrez, Ting Fang Tan, and Daniel Shu Wei Ting. 2023. Large language models in medicine. *Nature medicine* 29, 8 (2023), 1930–1940.

[17] Qian Wang et al. 2024. Epistemic Alignment: Matching User Epistemic Goals in LLM Interfaces. OpenReview. OpenReview submission on epistemic alignment frameworks for LLM interfaces.

[18] Zhepei Wei, Xiao Yang, Kai Sun, Jiaqi Wang, Rulin Shao, Sean Chen, Mohammad Kachuee, Teja Gollapudi, Tony Liao, Nicolas Scheffer, Rakesh Wanga, Anuj Kumar, Yu Meng, Wen tau Yih, and Xin Luna Dong. 2025. TruthRL: Incentivizing Truthful LLMs via Reinforcement Learning. arXiv:2509.25760 [cs.CL] https://arxiv.org/abs/2509.25760

[19] Thomas Winninger, Boussad Addad, and Katarzyna Kapusta. 2025. Using Mechanistic Interpretability to Craft Adversarial Attacks against Large Language Models. arXiv:2503.06269 [cs.LG] https://arxiv.org/abs/2503.06269

[20] Ziwei Xu, Sanjay Jain, and Mohan Kankanhalli. 2024. Hallucination is Inevitable: An Innate Limitation of Large Language Models. arXiv:2401.11817 [cs.CL]

[21] Alex L. Zhang, Tim Kraska, and Omar Khattab. 2025. Recursive Language Models. arXiv:2512.24601 [cs.AI] https://arxiv.org/abs/2512.24601